



Recursive fine-tuning damages a register, not a corpus and not the model globally

GPT-2 345M / top-k



SmolLM2 1.7B / top-k



Ratio: wiki flat, far-OOD declining

